

Course reminders

- Homework 1 due tonight
- Start forming final project groups & ideas (survey due Weds April 22)

Actor Critic Methods

CS 224R

1

Recap of Last Time

Online reinforcement learning with policy gradients

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \left(\left(\sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) \right) - b \right)$$

samples from policy policy log likelihood reward to go baseline

Do more of the above average stuff, less of the below average stuff.

Other notes about policy gradients: Need to collect entire trajectory before updating.
Does not rely on Markov property -> can use observations

3

2

Recap of Last Time

Vanilla policy gradients is fully **on-policy**.

—> only one gradient step per batch of data

Importance weights for off-policy policy gradient:

$$\nabla_{\theta'} J(\theta') \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \frac{\pi_{\theta'}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t})}{\pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t})} \nabla_{\theta'} \log \pi_{\theta'}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \left(\left(\sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) \right) - b \right)$$

Full algorithm:

1. sample $\{\tau^i\}$ from $\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$
2. compute $\nabla_{\theta} J(\theta)$ using $\{\tau^i\}$
3. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

Can take **multiple gradient steps** on the same batch

-> still very close to on-policy

attribute of online RL methods

Definitions.

on-policy: update uses only data from current policy

off-policy: update can reuse data from other, past policies

4

The plan for today

Actor critic methods

1. Improving policy gradients
 - a. Sample & directly supervise
 - b. Use your own estimate
 - c. Something in-between?
3. Off-policy actor-critic
 - a. Importance weights & constraining step size
 - b. Full off-policy version with replay buffers

Key learning goals:

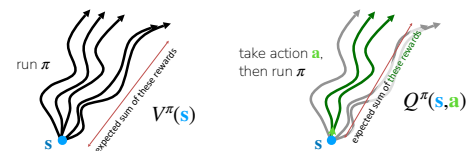
- How to estimate how good a state and action is for a policy
- How to use those estimates to form a more efficient RL algorithm

5

Revisiting Some Useful Objects

value function $V^{\pi}(\mathbf{s})$ - future expected rewards starting at $\mathbf{s}$ and following π

Q-function $Q^{\pi}(\mathbf{s}, \mathbf{a})$ - future expected rewards starting at $\mathbf{s}$, taking $\mathbf{a}$, then following π



$$\text{Useful relation: } V^{\pi}(\mathbf{s}) = \mathbb{E}_{\mathbf{a} \sim \pi(\cdot | \mathbf{s})} [Q^{\pi}(\mathbf{s}, \mathbf{a})]$$

advantage $A^{\pi}(\mathbf{s}, \mathbf{a})$ - how much better it is to take $\mathbf{a}$ than to follow policy π at state $\mathbf{s}$

$$A^{\pi}(\mathbf{s}, \mathbf{a}) = Q^{\pi}(\mathbf{s}, \mathbf{a}) - V^{\pi}(\mathbf{s})$$

Let's look at an example

Reward = 1 if I can play it in a month, 0 otherwise

Question:
For each action, what are these?

Value function: $V^\pi(s_t) = ?$
Q function: $Q^\pi(s_t, a_t) = ?$
Advantage function: $A^\pi(s_t, a_t) = ?$

Current policy: $\pi(a_1 | s) = 1 \forall s$

What is dissatisfying about policy gradients?

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) \left(\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) - b \right)$$

samples from policy policy log likelihood reward to go baseline

τ^4 : one small step forward then falls backwards

τ^2 : folds only the sleeves
 τ^3 : flattens the jacket but does not fold it
 τ^4 : folds the jacket

-> pushes down likelihood of step forward -> with sparse rewards, don't utilize τ^2 or τ^3

Policy gradients doesn't make efficient use of data! Can we learn what is good & bad?

Improving policy gradients

Estimating reward to go

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) \underbrace{\left(\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) \right)}_{\text{"reward to go"}}$$

estimate of future rewards if we take $a_{i,t}$ in state $s_{i,t}$

Can we get a better estimate?

$$\sum_{t'=t}^T p(s_{i,t'}, a_{i,t'}) \rightarrow \sum_{t'=t}^T E_{\pi_\theta} [r(s_{i,t'}, a_{i,t'}) | s_{i,t}, a_{i,t}] = Q(s_{i,t}, a_{i,t})$$

true expected rewards to go This would be way better!

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) Q(s_{i,t}, a_{i,t})$$

Improving policy gradients

What about baselines?

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) (Q(s_{i,t}, a_{i,t}) - \cancel{b})$$

b = average reward = $\frac{1}{N} \sum_i Q(s_{i,t}, a_{i,t})$

$$V(s_t) = E_{a_t \sim \pi_\theta(\cdot | s_t)} [Q(s_t, a_t)]$$

Recall: $A^\pi(s_t, a_t) = Q^\pi(s_t, a_t) - V^\pi(s_t)$

With baseline:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) A^\pi(s_{i,t}, a_{i,t})$$

Better estimates of A^π lead to less noisy gradients!

Online RL Outline

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) A^\pi(s_{i,t}, a_{i,t})$$

First: Initialize the policy (randomly, with imitation learning, with heuristics)

Run policy to collect batch of data

Fit model to estimate expected return

Improve policy

Online RL Outline

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) A^\pi(s_{i,t}, a_{i,t})$$

First: Initialize the policy (randomly, with imitation learning, with heuristics)

Run policy to collect batch of data

Fit model to estimate expected return

Estimate V^π, Q^π , or A^π

Improve policy

$\theta \leftarrow \theta + \nabla_\theta J(\theta)$

The plan for today

Actor critic methods

1. Improving policy gradients
2. **How to estimate the value of a policy**
 - a. Sample & directly supervise
 - b. Use your own estimate
 - c. Something in-between?
3. Off-policy actor-critic
 - a. Importance weights & constraining step size
 - b. Full off-policy version with replay buffers

Key learning goals:

- How to estimate how good a state and action is for a policy
- How to use those estimates to form a better RL algorithm

13

Estimating expected return


$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) A^{\pi}(\mathbf{s}_{i,t}, \mathbf{a}_{i,t})$$

Should we fit V^{π} , Q^{π} , or A^{π} ?

$$\begin{aligned} A^{\pi}(\mathbf{s}_t, \mathbf{a}_t) &= Q^{\pi}(\mathbf{s}_t, \mathbf{a}_t) - \hat{V}^{\pi}(\mathbf{s}_t) \\ Q^{\pi}(\mathbf{s}_t, \mathbf{a}_t) &= \sum_{t'=t}^T E_{\pi_{\theta}} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t, \mathbf{a}_t] \\ &= r(\mathbf{s}_t, \mathbf{a}_t) + \sum_{t'=t+1}^T E_{\pi_{\theta}} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t, \mathbf{a}_t] \\ &= r(\mathbf{s}_t, \mathbf{a}_t) + E_{\pi_{\theta}} [V^{\pi}(\mathbf{s}_{t+1}) | \mathbf{s}_t, \mathbf{a}_t] \\ &\approx r(\mathbf{s}_t, \mathbf{a}_t) + V^{\pi}(\mathbf{s}_{t+1}) \quad (\text{use the sampled } \mathbf{s}_{t+1}) \\ A^{\pi}(\mathbf{s}_t, \mathbf{a}_t) &\approx r(\mathbf{s}_t, \mathbf{a}_t) + V^{\pi}(\mathbf{s}_{t+1}) - \hat{V}^{\pi}(\mathbf{s}_t) \end{aligned}$$

Let's just fit V^{π} !

14




parameters ϕ

Estimating V^{π}

Version 1: Monte Carlo estimation

$$V^{\pi}(\mathbf{s}_t) \approx \sum_{t'=t}^T r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \quad \leftarrow \text{original single sample estimate}$$




15

Estimating V^{π}

Version 1: Monte Carlo estimation

$$V^{\pi}(\mathbf{s}_t) \approx \sum_{t'=t}^T r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \quad \leftarrow \text{original single sample estimate}$$

$$V^{\pi}(\mathbf{s}_t) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) \quad \leftarrow \text{multi-sample estimate (but can't reset the world)}$$



16

Estimating V^{π}

Version 1: Monte Carlo estimation

$$V^{\pi}(\mathbf{s}_t) \approx \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) \quad \leftarrow \text{original single sample estimate}$$

$$V^{\pi}(\mathbf{s}_t) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) \quad \leftarrow \text{multi-sample estimate (but can't reset the world)}$$

Step 1: Aggregate dataset of single sample estimates:


$$\left\{ \left(\mathbf{s}_{i,t}, \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) \right) \right\}$$

$y_{i,t}$

Step 2: Supervised learning to fit estimated value function

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_i \left\| \hat{V}_{\phi}^{\pi}(\mathbf{s}_i) - y_i \right\|^2$$

17



Estimating V^{π}

Version 2: Bootstrapping

$$\text{Monte Carlo target: } y_{i,t} = \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'})$$

$$\begin{aligned} \text{ideal target: } y_{i,t} &= \sum_{t'=t}^T E_{\pi_{\theta}} [r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) | \mathbf{s}_{i,t}] \approx r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \sum_{t'=t+1}^T E_{\pi_{\theta}} [r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) | \mathbf{s}_{i,t+1}] \\ &\approx r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + V^{\pi}(\mathbf{s}_{i,t+1}) \approx r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \hat{V}_{\phi}^{\pi}(\mathbf{s}_{i,t+1}) \end{aligned}$$

directly use previous fitted value function!

Step 1: Aggregate dataset of "bootstrapped" estimates:


$$\text{training data: } \left\{ \left(\mathbf{s}_{i,t}, \underbrace{r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \hat{V}_{\phi}^{\pi}(\mathbf{s}_{i,t+1})}_{y_{i,t}} \right) \right\} \quad \leftarrow \text{update labels every gradient update!}$$

Step 2: Supervised learning to fit estimated value function

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_i \left\| \hat{V}_{\phi}^{\pi}(\mathbf{s}_i) - y_i \right\|^2$$

Also referred to as a form of **temporal difference learning**

18




parameters ϕ

Estimating V^π : Monte Carlo vs. Bootstrap

Monte Carlo $y_{i,t} = \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'})$ supervise with roll-out's summed rewards
Bootstrapped $y_{i,t} = r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \hat{V}_\phi^\pi(\mathbf{s}_{i,t+1})$ supervise using reward and current value estimate

Let's look at an example:




Think-pair-share:
 $\hat{V}_{MC}^\pi(\mathbf{s}) \approx ?$ $\hat{V}_{MC}^\pi(\mathbf{s}) \approx ?$
 $\hat{V}_{TD}^\pi(\mathbf{s}) \approx ?$ $\hat{V}_{TD}^\pi(\mathbf{s}) \approx ?$

Is there middle ground?
Can we balance bias and variance?

Estimating V^π : Monte Carlo vs. Bootstrap

Monte Carlo $y_{i,t} = \sum_{t'=t}^T r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'})$ supervise with roll-out's summed rewards
Bootstrapped $y_{i,t} = r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \hat{V}_\phi^\pi(\mathbf{s}_{i,t+1})$ supervise using reward and current value estimate
N-step returns $y_{i,t} = \sum_{t'=t}^{t+n-1} r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) + \hat{V}_\phi^\pi(\mathbf{s}_{i,t+n})$



-> less variance than MC
-> lower bias than 1-step bootstrap

$n > 1, n < T$ often works the best in practice!

Aside: discount factors

$y_{i,t} \approx r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \hat{V}_\phi^\pi(\mathbf{s}_{i,t+1})$


$\mathcal{L}(\phi) = \frac{1}{2} \sum_i \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$

what if T (episode length) is ∞ ?
 $\hat{V}_\phi^\pi$ can get infinitely large in many cases

simple trick: better to get rewards sooner than later γ changes the MDP:



$y_{i,t} \approx r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) + \gamma \hat{V}_\phi^\pi(\mathbf{s}_{i,t+1})$

discount factor $\gamma \in [0, 1]$ (0.99 works well)



A Full Algorithm Walkthrough

1. Sample batch of data $\{(\mathbf{s}_{1,i}, \mathbf{a}_{1,i}, \dots, \mathbf{s}_{T,i}, \mathbf{a}_{T,i})\}$ from π_θ
2. Fit $\hat{V}_\phi^{\pi_\theta}$ to summed rewards in data
3. Evaluate $\hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) = r(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) + \gamma \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t+1,i}) - \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t,i}) \quad \forall t, i$
4. Evaluate $\nabla_\theta J(\theta) \approx \sum_{t,i} \nabla_\theta \log \pi_\theta(\mathbf{a}_{t,i} | \mathbf{s}_{t,i}) \hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i})$
5. Update $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$



$y_{i,t} = \sum_{t'=t}^{t+n-1} \gamma^{t'-t} r(\mathbf{s}_{i,t'}, \mathbf{a}_{i,t'}) + \gamma^n \hat{V}_\phi^\pi(\mathbf{s}_{i,t+n})$
 $\mathcal{L}(\phi) = \frac{1}{2} \sum_i \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$

"actor-critic algorithm"


Review so far

Algorithms
Policy gradients: observe what is good vs. bad, then do more of good stuff
Actor-critic: learn to estimate what is good vs. bad, then do more of the good stuff

i.e. get a better policy gradient by using neural network to estimate value function

How to estimate value: "policy evaluation"

- Supervised learning directly on observed sum of future rewards
- Supervised learning on current reward + value estimate of next state
- Hybrid: Supervised learning on sum of next n rewards + value of state after that



What about off-policy versions?

The plan for today

- Actor critic methods**
1. Improving policy gradients
 2. How to estimate the value of a policy ("policy evaluation")
 - a. Sample & directly supervise ("Monte Carlo estimation")
 - b. Use your own estimate ("bootstrapping", "temporal difference learning")
 - c. Something in-between? ("N-step returns")
 3. **Off-policy actor-critic**
 - a. Importance weights & constraining step size
 - b. Full off-policy version with replay buffers

- Key learning goals:
- How to estimate how good a state and action is for a policy
 - How to use those estimates to form a better RL algorithm

Off-Policy Actor-Critic Methods

Version 1: Multiple Gradient Steps

- 1. Sample batch of data $\{(\mathbf{s}_{1,i}, \mathbf{a}_{1,i}, \dots, \mathbf{s}_{T,i}, \mathbf{a}_{T,i})\}$ from π_θ
- 2. Fit $\hat{V}_\phi^{\pi_\theta}$ to summed rewards in data
- 3. Evaluate $\hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) = r(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) + \gamma \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t+1,i}) - \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t,i}) \quad \forall t, i$
- 4. Evaluate $\nabla_\theta J(\theta) \approx \sum_{t,i} \nabla_\theta \log \pi_\theta(\mathbf{a}_{t,i} | \mathbf{s}_{t,i}) \hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i})$ <- use importance weights here
- 5. Update $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

25

Off-Policy Actor-Critic Methods

Version 1: Multiple Gradient Steps

- 1. Sample batch of data $\{(\mathbf{s}_{1,i}, \mathbf{a}_{1,i}, \dots, \mathbf{s}_{T,i}, \mathbf{a}_{T,i})\}$ from π_θ
 - 2. Fit $\hat{V}_\phi^{\pi_\theta}$ to summed rewards in data
 - 3. Evaluate $\hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) = r(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) + \gamma \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t+1,i}) - \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t,i}) \quad \forall t, i$
 - 4. Evaluate $\nabla_{\theta'} J(\theta') \approx \sum_{t,i} \frac{\pi_{\theta'}(\mathbf{a}_{t,i} | \mathbf{s}_{t,i})}{\pi_\theta(\mathbf{a}_{t,i} | \mathbf{s}_{t,i})} \nabla_{\theta'} \log \pi_{\theta'}(\mathbf{a}_{t,i} | \mathbf{s}_{t,i}) \hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i})$ <- use importance weights here
 - 5. Update $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
- Policy will increase probability on actions with high advantages
- Advantages based on old policy.

What can go wrong if you take too many gradient steps?

26

Off-Policy Actor-Critic Methods

Version 1: Multiple Gradient Steps

- 1. Sample batch of data $\{(\mathbf{s}_{1,i}, \mathbf{a}_{1,i}, \dots, \mathbf{s}_{T,i}, \mathbf{a}_{T,i})\}$ from π_θ
 - 2. Fit $\hat{V}_\phi^{\pi_\theta}$ to summed rewards in data
 - 3. Evaluate $\hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) = r(\mathbf{s}_{t,i}, \mathbf{a}_{t,i}) + \gamma \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t+1,i}) - \hat{V}_\phi^{\pi_\theta}(\mathbf{s}_{t,i}) \quad \forall t, i$
 - 4. Evaluate $\nabla_{\theta'} J(\theta') \approx \sum_{t,i} \frac{\pi_{\theta'}(\mathbf{a}_{t,i} | \mathbf{s}_{t,i})}{\pi_\theta(\mathbf{a}_{t,i} | \mathbf{s}_{t,i})} \nabla_{\theta'} \log \pi_{\theta'}(\mathbf{a}_{t,i} | \mathbf{s}_{t,i}) \hat{A}^{\pi_\theta}(\mathbf{s}_{t,i}, \mathbf{a}_{t,i})$ <- use importance weights here
 - 5. Update $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
- Policy will increase probability on actions with high advantages
- Advantages are based on old policy, will get out-dated

What can go wrong if you take too many gradient steps?

- Idea 1: Use KL constraint on policy.
We will see this in LLM preference optimization
 $\mathbb{E}_{\mathbf{s} \sim \pi_\theta} [D_{KL}(\pi_{\theta'}(\cdot | \mathbf{s}) \| \pi_\theta(\cdot | \mathbf{s}))] \leq \delta$
- Idea 2: Can we bound the importance weights?
Doesn't directly constrain policy, but removes incentives
-> Key idea behind proximal policy optimization (PPO)

27

Off-Policy Actor-Critic Methods

- So far:
 - use one batch of policy data for one gradient step (fully on-policy)
 - use one batch of policy data for multiple gradient steps (starting to be off-policy)

Can we be even more off-policy?

Can we reuse data from previous batches, i.e. all of the past trial-and-error data?


- Key ideas:
 - maintain a buffer of all past data "replay buffer"
 - adjust equations to remove on-policy assumptions

28

Off-Policy Actor-Critic Methods

Version 2: Replay buffers

- actor-critic algorithm:
- 1. collect experience $\{\mathbf{s}_i, \mathbf{a}_i\}$ from $\pi_\theta(\mathbf{a} | \mathbf{s})$ } Add this to replay buffer
 - 2. fit $\hat{V}_\phi^\pi(\mathbf{s})$ to sampled reward sums
 - 3. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
 - 4. $\nabla_\theta J(\theta) \approx \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i | \mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
 - 5. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
- Do this on minibatch sampled from all previous data



Off-Policy Actor-Critic Methods

Version 2: Replay buffers

- online actor-critic algorithm:
- 1. collect experience $\{\mathbf{s}_i, \mathbf{a}_i\}$ from $\pi_\theta(\mathbf{a} | \mathbf{s})$ (& add to replay buffer)
 - 2. sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}$ from buffer $\mathcal{R}$
 - 3. update $\hat{V}_\phi^\pi$ using targets $y_i = r_i + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$ to each $\mathbf{s}_i$
 - 4. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
 - 5. $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i | \mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
 - 6. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
- not the right target value
- not the action π_θ would have taken!

replay buffer

mini batch size

$\mathcal{L}(\phi) = \frac{1}{N} \sum_i \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$

The algorithm is currently broken 😞

29

Slide adapted from Sergey Levine

30

Version 2: Fixing the value function

➡ 1. take action $\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s})$, get $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$, store in $\mathcal{R}$

- sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}$ from buffer $\mathcal{R}$
- update $\hat{V}_\phi^\pi$ using targets $y_i \leftarrow r_i + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$ to each $\mathbf{s}_i$
- evaluate $A^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
- $\nabla_\theta J(\theta) \leftarrow \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i | \mathbf{s}_i) A^\pi(\mathbf{s}_i, \mathbf{a}_i)$
- $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

not the right target value

where does this come from?

$$Q^\pi(s_t) = \sum_{t'=t}^T E_{\pi_\theta} [r(s_{t'}, \mathbf{a}_{t'}) | s_t]$$

“total reward we get if we take $\mathbf{a}_t$ in $\mathbf{s}_t$...
... and then follow the policy π ”

3. update $\hat{Q}_\phi^\pi$ using targets $y_i = r_i + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$ for each $\mathbf{s}_i, \mathbf{a}_i$
 $= r_i + \gamma \hat{Q}_\phi^\pi(\mathbf{s}'_i, \mathbf{a}'_i)$
not from replay buffer $\mathcal{R}$!
 $\mathbf{a}'_i \sim \pi_\theta(\mathbf{a}'_i | \mathbf{s}'_i)$

$$\mathcal{L}(\phi) = \frac{1}{N} \sum_i \left\| \hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a}_i) - y_i \right\|^2$$

$$V^\pi(\mathbf{s}_t) = \sum_{t'=t}^T E_{\pi_\theta} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t] = E_{\mathbf{a} \sim \pi(\mathbf{a}_t | \mathbf{s}_t)} [Q(\mathbf{s}_t, \mathbf{a}_t)]$$

31
Slide adapted from Se

Slide adapted from Sergey Levine Slide adapted from Sergey Levine

Version 2: Fixing the policy update

➡ 1. take action $\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s})$, get $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$, store in $\mathcal{R}$

2. sample batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}$ from buffer $\mathcal{R}$
3. update $\hat{Q}_\phi$ using targets $y_i = r_i + \gamma \hat{Q}_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$ for each $\mathbf{s}_i, \mathbf{a}_i$
4. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = Q(\mathbf{s}_i, \mathbf{a}_i) - \hat{V}_\phi(\mathbf{s}_i)$
5. $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum \nabla_{\theta} \log \pi_{\theta}(\mathbf{s}_i, \mathbf{a}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
6. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

not the action π_θ would have taken!

use the same trick, but this time for $\mathbf{a}_i$ rather than $\mathbf{a}_i'$!

sample $\mathbf{a}_i^\pi \sim \pi_\theta(\mathbf{a}|\mathbf{s}_i)$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_i^{\pi} | \mathbf{s}_i) \hat{A}^{\pi}(\mathbf{s}_i, \mathbf{a}_i^{\pi})$$

not from replay buffer $\mathcal{R}$! $\uparrow$ higher variance, but convenient / why is higher variance OK here

in practice: $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_i^{\pi} | \mathbf{s}_i) \hat{Q}^{\pi}(\mathbf{s}_i, \mathbf{a}_i^{\pi})$

32

Version 2: Anything else?

➡ 1. take action $\mathbf{a} \sim \pi_{\theta}(\mathbf{a}|\mathbf{s})$, get $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$, store in $\mathcal{R}$

- sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, \mathbf{r}_i, \mathbf{s}'_i\}$ from buffer $\mathcal{R}$
- update $\hat{Q}^\pi_\theta$ using targets $y_i = r_i + \gamma \hat{Q}^\pi_\theta(\mathbf{s}'_i, \mathbf{a}'_i)$ for each $\mathbf{s}_i, \mathbf{a}_i$
- $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}^\pi_i | \mathbf{s}_i) \hat{Q}^\pi_\theta(\mathbf{s}_i, \mathbf{a}^\pi_i)$ where $\mathbf{a}^\pi_i \sim \pi_\theta(\mathbf{a} | \mathbf{s}_i)$
- $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Any remaining problems?

$\mathbf{s}_i$ didn't come from $p_\theta(\mathbf{s})$

nothing we can do here, just accept it

intuition: we want optimal policy on $p_\theta(\mathbf{s})$

but we get optimal policy on a *broader* distribution

Slide adapted from Sergey Levine

Version 2: Some implementation details

➡ 1. take action $\mathbf{a} \sim \pi_{\theta}(\mathbf{a}|\mathbf{s})$, get $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$, store in $\mathcal{R}$

- sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}$ from buffer $\mathcal{R}$
- update $\hat{Q}^{\pi}_{\theta}$ using targets $y_i = r_i + \gamma \hat{Q}^{\pi}_{\theta}(\mathbf{s}'_i, \mathbf{a}'_i)$ for each $\mathbf{s}_i, \mathbf{a}_i$
- $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}^{\pi}_i | \mathbf{s}_i) \hat{Q}^{\pi}_{\theta}(\mathbf{s}_i, \mathbf{a}^{\pi}_i)$ where $\mathbf{a}^{\pi}_i \sim \pi_{\theta}(\mathbf{a} | \mathbf{s}_i)$
- $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

also fancier ways to fit Q-functions
(more on this in next two lectures)


can also use **reparameterization trick** to better estimate the gradient (for Gaussian policy)

Example practical algorithm:

Haarnoja, Zhou, Abbeel, Levine. Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor. 2018.

34

The plan for today



- + Off-policy with replay buffer (e.g. soft actor-critic) can be far more data efficient

- They can also generally be harder to tune hyperparameters, less stable (than e.g. PPO)

actor critic methods

1. Improving policy gradients
2. How to estimate the value of a policy ("**policy evaluation**")
 - a. Sample & directly supervise ("**Monte Carlo estimation**")
 - b. Use your own estimate ("**bootstrapping**", "**temporal difference learning**")
 - c. Something in-between? ("**N-step returns**")
3. Off-policy actor-critic
 - a. Importance weights & constraining step size
 - b. Full off-policy version with replay buffers

Key learning goals:

- How to estimate how good a state and action is for a policy
- How to use those estimates to form a better RL algorithm

35

36

Next Week

Q-learning (last big online RL method)

+

Practical implementation of online RL algorithms

Course reminders

- Homework 1 due tonight
- Form final project groups & ideas (survey due Weds April 22)